

A wide-angle photograph of the San Francisco skyline at sunset. The sky is a mix of orange, pink, and blue, with scattered clouds. The city's buildings are silhouetted against the bright sky, with some windows glowing with light. The Transamerica Pyramid is a prominent feature on the left. In the foreground, the city's dense residential areas are visible, and the water of the bay is calm with a few boats.

Price My Rental

Apartment Price Recommendation Engine

[Get Started](#)

Taken from actual
craigslist listing on
6/15/15

Listed price on craigslist

Submit

Web app - PriceMyRental.io

Your model-recommended price is: \$3782

Your initial stated price was: \$3950

About:

Price My Rental predicts prices for rental listings in San Francisco. For landlords, the app can recommend a fair monthly price for your property. For tenants, it can help determine if you're overpaying or getting a good deal.

For more information on how it works, please visit the [GitHub repo](#)

 [Contact Me](#)

[Source on Github](#)

Copyright © Jon Oleson 2015

Web app - PriceMyRental.io

What is a rental worth?


- San Francisco has one of the hottest rental markets in the US, and a ton of variation
- I looked at 9 months of SF craigslist housing rental data, and built a model that accurately predicts market prices for units

Techniques


- KD-Trees: Finding nearest-neighbors price medians
 - Space-partitioning model that splits data on medians, allowing the user to efficiently find approximate nearest neighbors
 -
- TF-IDF + NMF: Discovering latent features in listing descriptions
 - TF-IDF: Numerical statistic representing how important a word is to a document within a larger corpus
 - NMF: Factorizing the TF-IDF matrix into component matrices, one of which represents how highly each listing description ranks on each latent feature we're looking to uncover (in our case, 4 features)

Techniques

kD-Trees




Split longer dimension near data media




term-doc matrix

=



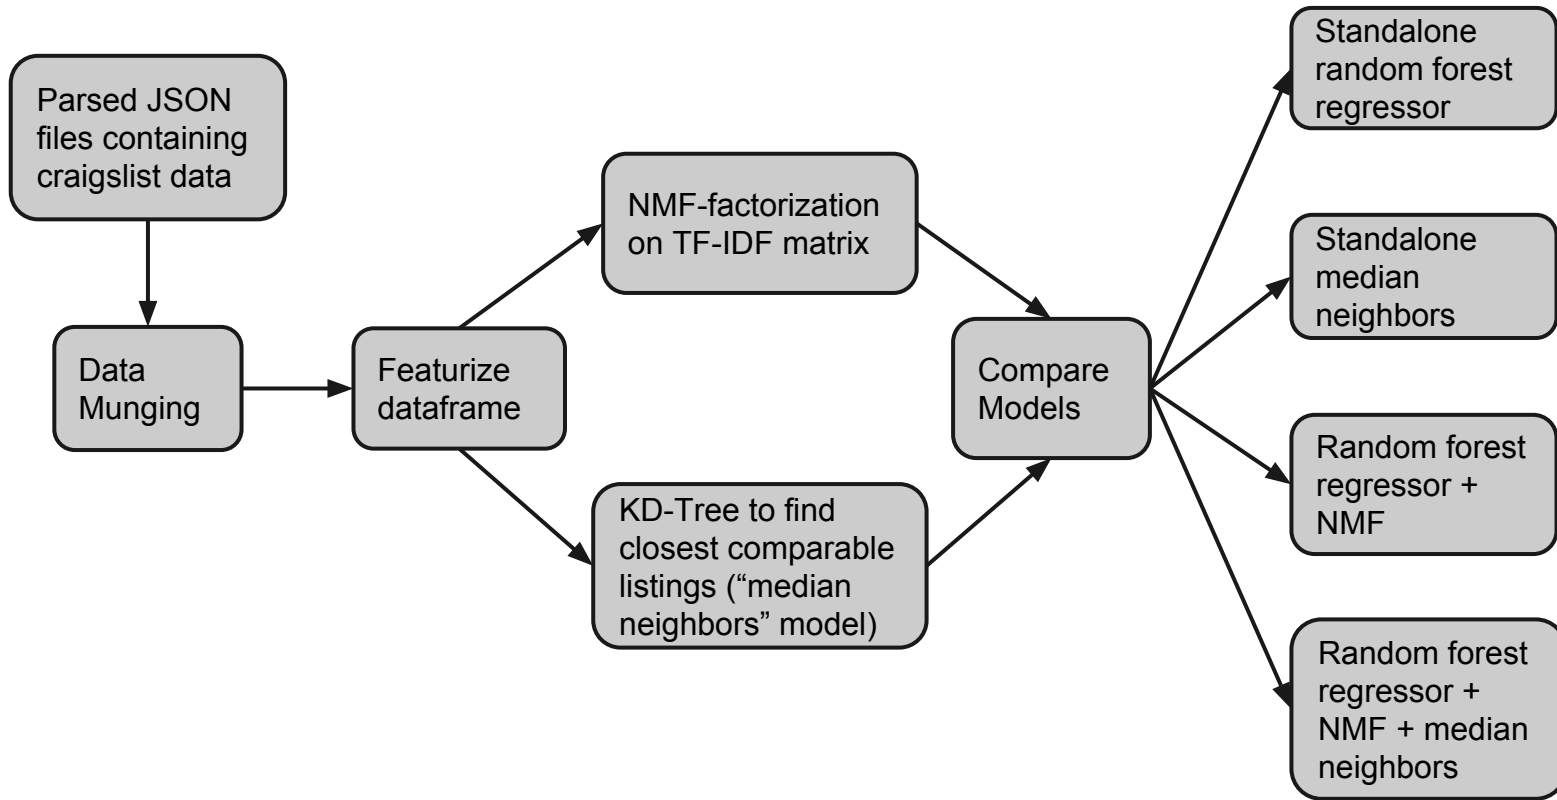
term-topic matrix

×

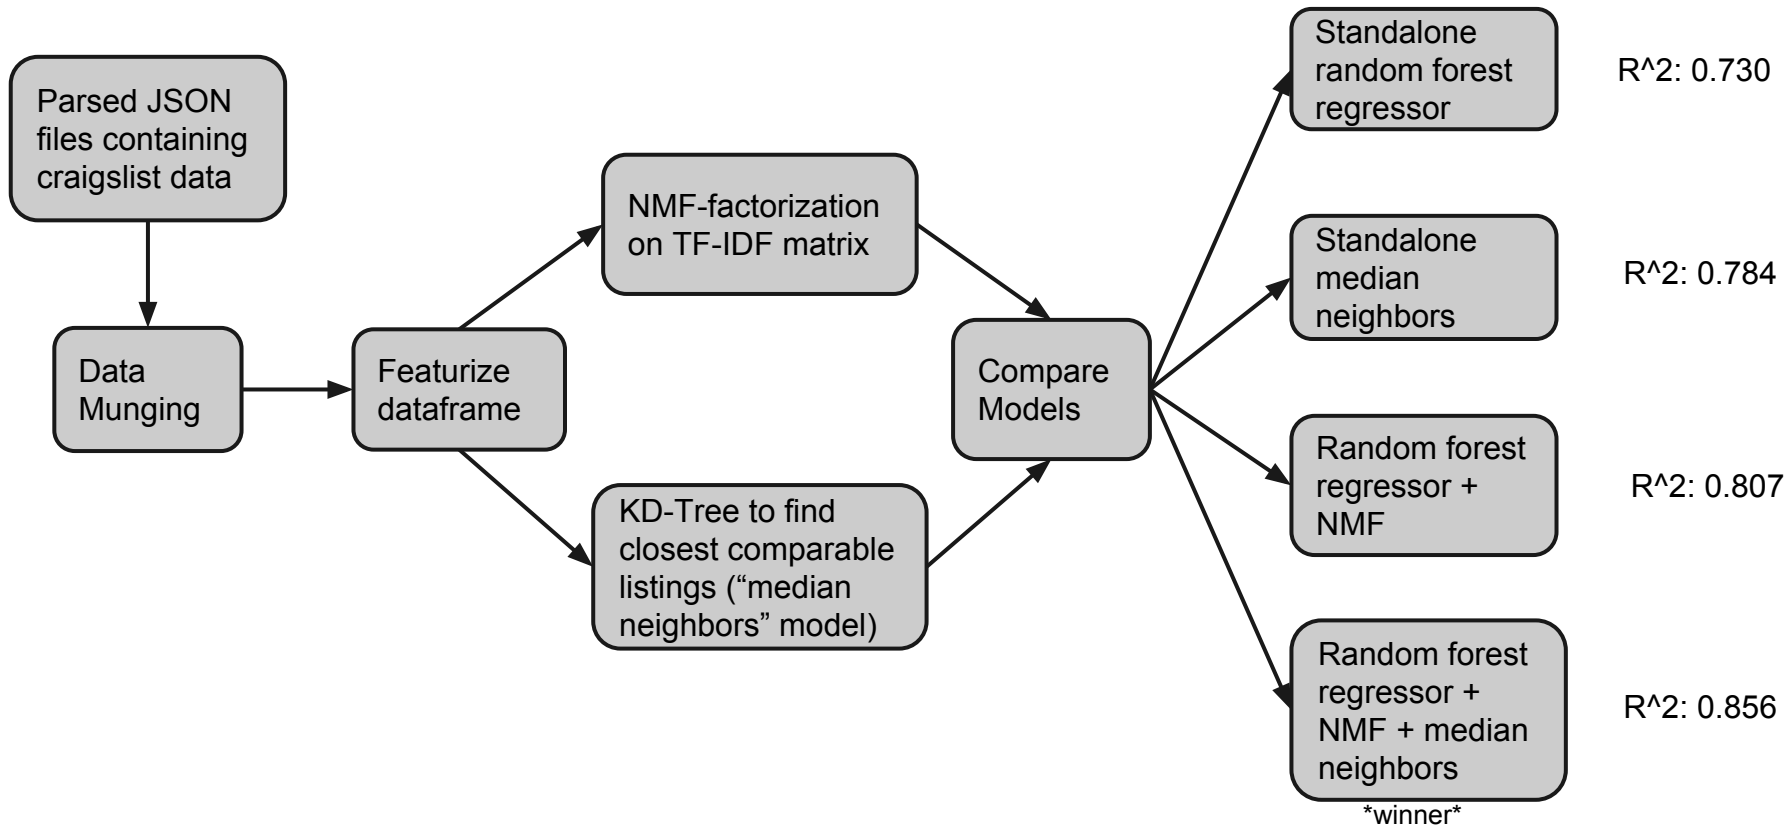


topic-doc matrix

NMF



Data Pipeline + Model Comparison



Data Pipeline + Model Comparison

Future Work

- Build functionality for more cities
- More rigorously test performance of different regression models
- Fix occasional discontinuities in model predictions (e.g., price decreasing when bedrooms increased for a given listing, holding all else equal)
- Add flexibility to prevent failure with edge cases (e.g., listing with 5 beds and 1 bath)

Libraries Used

- Scikit-Learn
- Pandas
- Numpy
- NLTK
- Geopy
- Flask
- Datetime
- urllib2
- cPickle